


Comp Org Overview

- 4 subsystems
 - Processor-core
 - Signal Chain
 - Memory
 - Comms and UI

Signal Chain

i) ADC - DAC (+filter)

- myquist: sample freq $\geq 2^*$ (signal freq)

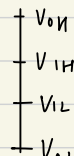
ii) Interfaces

- serial / parallel bits + single / bi direction
- input / output / I & O
- compatible if electrically + communication protocol

$V_{OH} \geq V_{IH}, V_{OL} \leq V_{IL}$

o/i : output/input ; H/L : high/low

min max



- differential signals have better noise tolerance and can be clocked at higher frequencies
 - greater margin
 - external interference cancellation
 - difference b/w V^+ & V^- remains constant
- communications protocol: no of bits, type of synchronization, types of data + formatting

- Parallel Data transfer

- synchronous
- preferred when sustained high speed is needed
- prone to signal skew + crosstalk ← limit max clock freq
 - ↓ frequency, number of lines (ground plane helps)
 - temp, capacitance, printed circuit board dimensions
- needs only strobe signal, simpler hw interface

- Serial data transfer

- more complex hw interface, slower
- preferred for long distances
- simplex (1D), half duplex (2-D but 1 at a time), full duplex
- sync (common clock signal); async (both have clock signals, not the same, pre-fixed clock freq is used)
- sync serial data transfer
 - master/slave · master receives data on posedge, passes on
 - ↳ provides clock / negedge, MSB first;
- SPI : serial peripheral interface : synchronous
 - > slave select low ⇒ start transfer
 - > MOSI / MISO used depending on direction of transfer
 - > multiple slaves allowed
- async serial data transfer
 - sync words are used,
 - possible skew as 2 separate clocks

- UART : universal asynchronous receiver transmitter
 - transfer {
 - > idle : logic 1
 - > start : logic 0
 - > baud rate : no of signal state changes / second
 - > parity bit
 - > stop : logic 1
 - receiver {
 - > LSB first : waveform will have MSB first
 - > internal clock runs at $k \times$ baud rate ($k=2, 4, 16$) so that sampling is at midpoint of data bit

Data transfer Mechanism

i) Polling / Programmed I/O

- CPU uses 100% of resources to poll I/O port continuously once control reg for data transfer have been initialised.
- inefficient use of CPU ; very simple SW implementation

ii) Interrupt

- signal (Interrupt request) $\xrightarrow{\text{IRQ}}$ CPU $\xrightarrow{\text{interrupt latency delay}}$ Interrupt service routine (ISR)
- uses a vector table w/ indexes corresponding to IR's and a col w/ ISRs
- complex hw, time efficient, allows prioritization

iii) DMA (Direct Memory Access)

- controls Data transfer, shares same bus as CPU
- mem - mem, I/O per - I/O per, mem - I/O per
- gen address, read write ops
- fetches data to DMAc buffer, then deposits it ($\text{src} \rightarrow \text{DMAC} \rightarrow \text{dest}$)
- issues DMA Interrupt to CPU after data transfer to inform that data is ready to be processed.
- good for specific formats

- Burst mode

- › transfers multiple units, suspends CPU for a while, fast

- Cycle stealing

- › one unit at a time can occur b/w CPU instructions/pipeline stages

- › slow, CPU is freeer, preferred in security systems

- › transparent mode

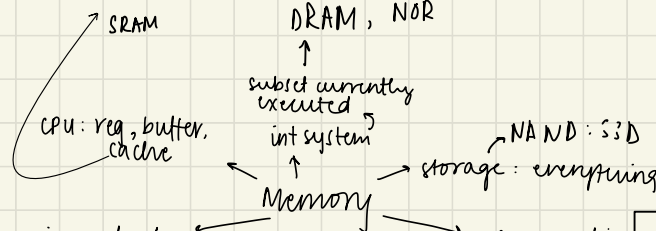
- transfer only when CPU is not using bus

- slowest

- least CPU disruptions

- complex hw

backup storage: HDD not so cost, ∞ read write



- faster, more expensive
- large: 6 transistors
 - ↳ 2: data flow
 - ↳ 4: store data
- low power consumption
- power supply \Rightarrow data stored
- can access specific locations

- slower, cheaper
- smaller: 1 transistor: charge flow
- 1 capacitor: holds charge
- 1 \rightarrow charged; 0 \rightarrow discharge
- every read destroys data: rewrite is needed

- periodic refresh, large power consumption

semi-conductor
SSD, SRAM, DRAM, Flash

volatile

SRAM } XIP
DRAM }

not exactly but they make it work

SSD

non volatile

EEPROM }
Flash }

NAND
NOR

- all cells $> V_T \Rightarrow$ bit line off = 0
- no XIP
- memory is accessed at page level
- less wires, used in mass storage
- combined I/O databus

- any cell has $V > V_T$ (logic 0) \Rightarrow bit line o/p = 0
- expensive as comp to NAND

- supports XIP
- can store program code
- random reading \checkmark
- erase at block lvl
- separate data, address lines

- flash
- limited read + write
- expensive
- longevity:
 - i) wear levelling
 - ii) external RAM buffer
 - iii) error correction code
- small & light as comp. to HDD
and faster

- logic \equiv orientation of magnetic elements
- \bigcirc platter \bigcirc track \curvearrowright sector \odot head \cdot no of surface
- cylinder track/surface
- $T_A = T_S + T_R$ (request to head = request to track + track to sector)
- T_T : after head in position to transfer data
- $T_R = \frac{0.5}{RPS}$; $T_T = \frac{N}{D_s \cdot D_T} \times \frac{1}{RPS}$ (D_T : sector/track, D_s : byte/sector)
↳ no of track to access
- disk defragmentation
- zone bit recording format
- LBA: logical block addressing (2^{30} bytes of data) can be addressed
- CHS: cylinder, head, sector
- \uparrow are addressing schemes
- performance: 100MB/s, 10000 RPM, 100000 IOPS

- FGT: floating gate transistor
- mosfet
- 2 states: Erased \rightarrow logic 1
Programmed \rightarrow 0
 $\rightarrow T$: off
- to program: large +ve voltage
- check state: inject voltage b/w
 $V_1 < V < V_T \leftarrow$ increased threshold prog
- if ON \Rightarrow erased
- if OFF \Rightarrow programmed

Cache

- fast go-b/w between CPU & system
- small parts of memory are transferred to the cache based on:
 - Principle of locality
 - ↳ locality of space (sequential) (local)
 - ↳ locality of time (loop) (temporal)
- cache mapping : direct mapped
 - > main memory address $\xrightarrow{\text{conv}}$ cache add.

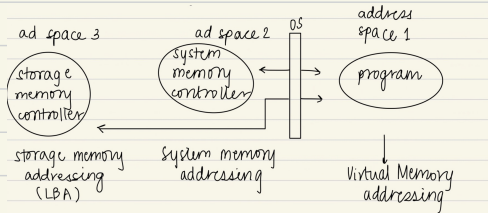
tag	block	offset
-----	-------	--------

 - ↳ same no of bits \uparrow
 - > offset bits = $\log_2 (\text{size of block})$ [size of cache block = size of MM block]
 - > block bits = $\log_2 (\text{size of MM} / (\text{size of cache} * \text{block size}))$
 - > tag bits = whatever is left.
- > tag values are compared to check if the memory is correct
- cache memory replacem^{en}t
 - LRU (least recently used) (slows down cache)
 - FIFO (one that has been there longest) (doesn't take runability into account)
- write policy
 - bus crowding system
 - i) write through : CPU writes in cache & main memory
 - ii) write back : memory is only updated when block is selected for replacement
 - ↑ cache coherency issue

- write miss : when block that CPU wants to update is not in the cache
 - i) write allocate : get block to cache
 - ii) write no-allocate : write directly in main memory
- effective access time : $EAT = H * \text{Access}_{\text{cache}} + (1-H) * (\text{Access}_{\text{cache}} + \text{Access}_{\text{mem}})$

Virtual Memory

- address translation is done by the OS



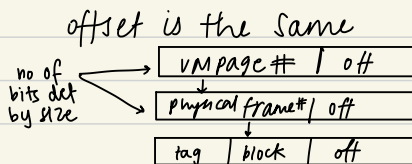
- Virtual address : address used by prog, generated by compiler
- VM can be contiguous
- address mapping : paging

VM size > memory size → the word size is same

VM	P/M/S/M	valid
page	frame	1
		0

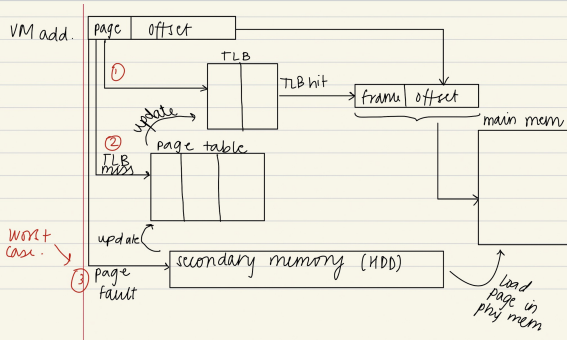
page table

↳ in MM



- ① VM → MM add
- ② CPU → cache hit/miss
- ③ hit → CPU takes info from cache
miss → cache manages transfers info from system/storage

TLB : translation lookaside buffer \rightarrow in processor
 \rightarrow SRAM
 \downarrow
subset of page table



w12

- unsigned number : $\text{carry} = 1 \Rightarrow \text{overflow}$
 $\text{overflow} \Rightarrow \text{nothing}$
 - signed number : $\text{overflow} = 1 \Rightarrow \text{overflow}$
 $\text{carry} \Rightarrow \text{not indicative of much}$
- } error

multibit arith : $\text{ADD } R0, R0, R1$ $\text{ADC } R2, R2, R3$

$\begin{array}{|c|c|} \hline R2 & R1 \\ \hline \end{array}$
 $\begin{array}{|c|c|} \hline R3 & R0 \\ \hline \end{array}$
 $\begin{array}{|c|c|} \hline R2 & R0 \\ \hline \end{array}$

↑
carry bit added

- 2's comp range : $-2^{N-1} \rightarrow 2^{N-1} - 1$

- fixed-point representation : range vs precision
 \downarrow
resolution

- floating point :

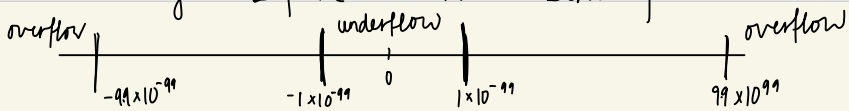
±		.				±	
---	--	---	--	--	--	---	--
- \uparrow sign $\underbrace{\hspace{2cm}}$ mantissa $\underbrace{\hspace{1cm}}$ exponent
base val \searrow position of radix

range : $(-9.99 \times 10^{99} \rightarrow 9.99 \times 10^{99})$

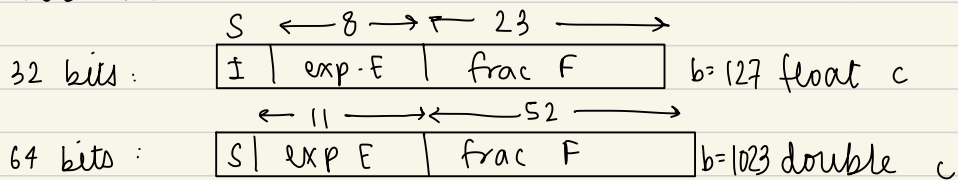
size of exp : range

size mantissa & value of exponent : precision

std : 1 digit before radix (binary : 1)



1 E E E 7 5 4



S: 0 \rightarrow +ve
1 \rightarrow -ve

$$\text{Exponent} = E - \text{Bias} (E - b)$$

F: $1 \cdot F$
 \hookrightarrow only frac starting from $2^{-1} \dots 2^{-23}/52$

$$(1)^S \times (1 + f_1 2^{-1} + f_2 2^{-2} + \dots) \times 2^{E-b}$$

$$\text{mantissa} = 1 \cdot f$$

E is a positive value, bias allows it to be signed.

8 bits $\Rightarrow 0 \rightarrow 255$
 $-127 \rightarrow 128$

?

normalised :

$$E: 1 \rightarrow 254$$

1. F : anything

denormalised

F : 00000000

O. F : nonzero

zero

F : 000000000

F : 00000, --- 000

infinity

E : 11011111

F: 0000 -- .000

NAN

$t : 1111111111$

f : non-zero

$$E = -126 \leftarrow$$

$$\text{not } -127$$

Handwritten: $\frac{1}{2}$ and $\frac{1}{2}$

32 bits

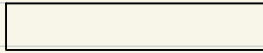
IEEE 754 (floating pt) \rightarrow precision changes

- range : $-2^{128} \rightarrow \sim 2^{128}$
- precision : less than 2^{-126}

fixed point \rightarrow uniform precision

- range $\approx 2^{32}$ unsigned
- max precision $\approx 2^{-32}$

sols.



offset : $2^{10} \therefore 10 \text{ bits}$

0 1100000000
1
11

$$\text{digital} = \frac{V}{V_{\text{ref}}} \times 2^8$$

~~25~~ 25

00100000
32

$$V = \frac{3}{8} = 0.75$$

$$O/P = I/P \times \frac{2^n}{V_{\text{ref}}}$$

$$\begin{matrix} 2.5 \\ 10.1 \\ 1.01 \times 2^1 \end{matrix}$$

$$\begin{matrix} I = E - 127 \\ E = 128 \end{matrix}$$